

Sorting Algorithms

```
/*
These functions are equivalent to std::sort(), taking
RandomAccessIterators
as a range [lo, hi) to be sorted. Elements between lo and hi
(including the
element pointed to by lo but excluding the element pointed to by
hi) will be
sorted into ascending order after the function call. Optionally, a
comparison
function object specifying a strict weak ordering may be specified
to replace
the default operator <.
```

```
These functions are not meant to compete with standard library
implementations
in terms of speed. Instead, they are meant to demonstrate
how common
sorting algorithms can be concisely implemented in C++.
*/
```

```
#include <algorithm>
#include <functional>
#include <iterator>
#include <vector>
```

```
/*
Quicksort repeatedly selects a pivot and partitions the range so
that elements
comparing less than the pivot precede the pivot, and elements
comparing greater
or equal follow it. Divide and conquer is then applied to both
sides of the
pivot until the original range is sorted. Despite having a worst
case of  $O(n^2)$ ,
quicksort is often faster in practice than merge sort and
heapsort, which both
have a worst case time complexity of  $O(n \log n)$ .
```

```
The pivot chosen in this implementation is always a middle element
of the range
to be sorted. To reduce the likelihood of encountering the worst
case, the pivot
can be chosen in better ways (e.g. randomly, or using the "median
of three"
technique).
```

```
Time Complexity (Average):  $O(n \log n)$ .
Time Complexity (Worst):  $O(n^2)$ .
Space Complexity:  $O(\log n)$  auxiliary stack space.
Stable?: No.
```

```

*/

template<class It, class Compare>
void quicksort(It lo, It hi, Compare comp) {
    if (hi - lo < 2) {
        return;
    }
    typedef typename std::iterator_traits<It>::value_type T;
    T pivot = *(lo + (hi - lo)/2);
    It i, j;
    for (i = lo, j = hi - 1; ; ++i, --j) {
        while (comp(*i, pivot)) {
            ++i;
        }
        while (comp(pivot, *j)) {
            --j;
        }
        if (i >= j) {
            break;
        }
        std::swap(*i, *j);
    }
    quicksort(lo, i, comp);
    quicksort(i, hi, comp);
}

```

```

template<class It>
void quicksort(It lo, It hi) {
    typedef typename std::iterator_traits<It>::value_type T;
    quicksort(lo, hi, std::less<T>());
}

```

/*
Merge sort first divides a list into n sublists of one element each, then recursively merges the sublists into sorted order until only a single sorted sublist remains. Merge sort is a stable sort, meaning that it preserves the relative order of elements which compare equal by operator < or the custom comparator given.

An analogous function in the C++ standard library is `std::stable_sort()`, except that the implementation here requires sufficient memory to be available. When $O(n)$ auxiliary memory is not available, `std::stable_sort()` falls back to a time complexity of $O(n \log^2 n)$ whereas the implementation here will simply fail.

Time Complexity (Average): $O(n \log n)$.
 Time Complexity (Worst): $O(n \log n)$.
 Space Complexity: $O(\log n)$ auxiliary stack space and $O(n)$ auxiliary heap space.
 Stable?: Yes.
 */

```
template<class It, class Compare>
void mergesort(It lo, It hi, Compare comp) {
    if (hi - lo < 2) {
        return;
    }
    It mid = lo + (hi - lo - 1)/2, a = lo, c = mid + 1;
    mergesort(lo, mid + 1, comp);
    mergesort(mid + 1, hi, comp);
    typedef typename std::iterator_traits<It>::value_type T;
    std::vector<T> merged;
    while (a <= mid && c < hi) {
        merged.push_back(comp(*c, *a) ? *c++ : *a++);
    }
    if (a > mid) {
        for (It it = c; it < hi; ++it) {
            merged.push_back(*it);
        }
    } else {
        for (It it = a; it <= mid; ++it) {
            merged.push_back(*it);
        }
    }
    for (int i = 0; i < hi - lo; i++) {
        *(lo + i) = merged[i];
    }
}
```

```
template<class It>
void mergesort(It lo, It hi) {
    typedef typename std::iterator_traits<It>::value_type T;
    mergesort(lo, hi, std::less<T>());
}
```

/*
 Heapsort first rearranges an array to satisfy the max-heap property. Then, it repeatedly pops the max element of the heap (the left, unsorted subrange), moving it to the beginning of the right, sorted subrange until the entire range is sorted. Heapsort has a better worst case time complexity than quicksort and also a better space complexity than merge sort.

The C++ standard library equivalent is calling `std::make_heap(lo, hi)`, followed by `std::sort_heap(lo, hi)`.

Time Complexity (Average): $O(n \log n)$.

Time Complexity (Worst): $O(n \log n)$.

Space Complexity: $O(1)$ auxiliary.

Stable?: No.

*/

```
template<class It, class Compare>
void heapsort(It lo, It hi, Compare comp) {
    typename std::iterator_traits<It>::value_type tmp;
    It i = lo + (hi - lo)/2, j = hi, parent, child;
    for (;;) {
        if (i <= lo) {
            if (--j == lo) {
                return;
            }
            tmp = *j;
            *j = *lo;
        } else {
            tmp = *(--i);
        }
        parent = i;
        child = lo + 2*(i - lo) + 1;
        while (child < j) {
            if (child + 1 < j && comp(*child, *(child + 1))) {
                child++;
            }
            if (!comp(tmp, *child)) {
                break;
            }
            *parent = *child;
            parent = child;
            child = lo + 2*(parent - lo) + 1;
        }
        *(lo + (parent - lo)) = tmp;
    }
}
```

```
template<class It>
void heapsort(It lo, It hi) {
    typedef typename std::iterator_traits<It>::value_type T;
    heapsort(lo, hi, std::less<T>());
}
```

/*

Comb sort is an improved bubble sort. While bubble sort increments the gap

between swapped elements for every inner loop iteration, comb sort fixes the gap size in the inner loop, decreasing it by a particular shrink factor in every iteration of the outer loop. The shrink factor of 1.3 is empirically determined to be the most effective.

Even though the worst case time complexity is $O(n^2)$, a well chosen shrink factor ensures that the gap sizes are co-prime, in turn requiring astronomically large n to make the algorithm exceed $O(n \log n)$ steps. On random arrays, comb sort is only 2-3 times slower than merge sort. Its short code length relative to its good performance makes it a worthwhile algorithm to remember.

Time Complexity (Worst): $O(n^2)$.
 Space Complexity: $O(1)$ auxiliary.
 Stable?: No.
 */

```
template<class It, class Compare>
void combsort(It lo, It hi, Compare comp) {
    int gap = hi - lo;
    bool swapped = true;
    while (gap > 1 || swapped) {
        if (gap > 1) {
            gap = (int)((double)gap / 1.3);
        }
        swapped = false;
        for (It it = lo; it + gap < hi; ++it) {
            if (comp(*(it + gap), *it)) {
                std::swap(*it, *(it + gap));
                swapped = true;
            }
        }
    }
}
```

```
template<class It>
void combsort(It lo, It hi) {
    typedef typename std::iterator_traits<It>::value_type T;
    combsort(lo, hi, std::less<T>());
}
```

/*
 Radix sort is used to sort integer elements with a constant number of bits in

linear time. This implementation only works on ranges pointing to unsigned integer primitives. The elements in the input range do not strictly have to be unsigned types, as long as their values are nonnegative integers.

In this implementation, a power of two is chosen to be the base for the sort so that bitwise operations can be easily used to extract digits. This avoids the need to use modulo and exponentiation, which are much more expensive operations. In practice, it's been demonstrated that 2^8 is the best choice for sorting 32-bit integers (approximately 5 times faster than `std::sort()`, and typically 2-4 faster than radix sort using any other power of two chosen as the base).

Time Complexity: $O(n \cdot w)$ for n integers of w bits each.

Space Complexity: $O(n + w)$ auxiliary.

*/

```
template<class UnsignedIt>
void radix_sort(UnsignedIt lo, UnsignedIt hi) {
    if (hi - lo < 2) {
        return;
    }
    const int radix_bits = 8;
    const int radix_base = 1 << radix_bits; // e.g.  $2^8 = 256$ 
    const int radix_mask = radix_base - 1; // e.g.  $2^8 - 1 = 0xFF$ 
    int num_bits = 8 * sizeof(*lo); // 8 bits per byte
    typedef typename std::iterator_traits<UnsignedIt>::value_type T;
    T *buf = new T[hi - lo];
    for (int pos = 0; pos < num_bits; pos += radix_bits) {
        int count[radix_base] = {0};
        for (UnsignedIt it = lo; it != hi; ++it) {
            count[( *it >> pos) & radix_mask]++;
        }
        T *bucket[radix_base], *curr = buf;
        for (int i = 0; i < radix_base; curr += count[i++]) {
            bucket[i] = curr;
        }
        for (UnsignedIt it = lo; it != hi; ++it) {
            *bucket[( *it >> pos) & radix_mask]++ = *it;
        }
        std::copy(buf, buf + (hi - lo), lo);
    }
    delete[] buf;
}
```

/** Example Usage and Output:

mergesort() with default comparisons: 1.32 1.41 1.62 1.73 2.58
2.72 3.14 4.67

mergesort() with 'compare_as_ints()': 1.41 1.73 1.32 1.62 2.72
2.58 3.14 4.67

Sorting five million integers...

std::sort(): 0.355s

quicksort(): 0.426s

mergesort(): 1.263s

heapsort(): 1.093s

combsort(): 0.827s

radix_sort(): 0.076s

***/

```
#include <cassert>
#include <cstdlib>
#include <ctime>
#include <iomanip>
#include <iostream>
#include <vector>
using namespace std;
```

```
template<class It>
void print_range(It lo, It hi) {
    while (lo != hi) {
        cout << *lo++ << " ";
    }
    cout << endl;
}
```

```
template<class It>
bool sorted(It lo, It hi) {
    while (++lo != hi) {
        if (*lo < *(lo - 1)) {
            return false;
        }
    }
    return true;
}
```

```
bool compare_as_ints(double i, double j) {
    return (int)i < (int)j;
}
```

```
int main () {
    { // Can be used to sort arrays like std::sort().
        int a[] = {32, 71, 12, 45, 26, 80, 53, 33};
        quicksort(a, a + 8);
    }
```

```

    assert(sorted(a, a + 8));
}
{ // STL containers work too.
  int a[] = {32, 71, 12, 45, 26, 80, 53, 33};
  vector<int> v(a, a + 8);
  quicksort(v.begin(), v.end());
  assert(sorted(v.begin(), v.end()));
}
{ // Reverse iterators work as expected.
  int a[] = {32, 71, 12, 45, 26, 80, 53, 33};
  vector<int> v(a, a + 8);
  heapsort(v.rbegin(), v.rend());
  assert(sorted(v.rbegin(), v.rend()));
}
{ // We can sort doubles just as well.
  double a[] = {1.1, -5.0, 6.23, 4.123, 155.2};
  vector<double> v(a, a + 5);
  combsort(v.begin(), v.end());
  assert(sorted(v.begin(), v.end()));
}
{ // Must use radix_sort with unsigned values, but sorting in
reverse works!
  int a[] = {32, 71, 12, 45, 26, 80, 53, 33};
  vector<int> v(a, a + 8);
  radix_sort(v.rbegin(), v.rend());
  assert(sorted(v.rbegin(), v.rend()));
}

// Example from:
http://www.cplusplus.com/reference/algorithm/stable\_sort
double a[] = {3.14, 1.41, 2.72, 4.67, 1.73, 1.32, 1.62, 2.58};
{
  vector<double> v(a, a + 8);
  cout << "mergesort() with default comparisons: ";
  mergesort(v.begin(), v.end());
  print_range(v.begin(), v.end());
}
{
  vector<double> v(a, a + 8);
  cout << "mergesort() with 'compare_as_ints()': ";
  mergesort(v.begin(), v.end(), compare_as_ints);
  print_range(v.begin(), v.end());
}
cout << "-----" << endl;

vector<int> v, v2;
for (int i = 0; i < 5000000; i++) {
  v.push_back((rand() & 0x7fff) | ((rand() & 0x7fff) << 15));
}
v2 = v;
cout << "Sorting five million integers..." << endl;
cout.precision(3);

```



```

#define test(sort_function) { \
    clock_t start = clock(); \
    sort_function(v.begin(), v.end()); \
    double t = (double)(clock() - start) / CLOCKS_PER_SEC; \
    cout << setw(14) << left << #sort_function "(): "; \
    cout << fixed << t << "s" << endl; \
    assert(sorted(v.begin(), v.end())); \
    v = v2; \
}

test(std::sort);
test(quicksort);
test(mergesort);
test(heap sort);
test(combsort);
test(radix_sort);
return 0;
}

```

Array Rotation

/*
These functions are equivalent to `std::rotate()`, taking three
iterators `lo`, `mid`,
and `hi` (`lo <= mid <= hi`) to perform a left rotation on the range
`[lo, hi)`. After
the function call, `[lo, hi)` will comprise of the concatenation of
the elements
originally in `[mid, hi) + [lo, mid)`. That is, the range `[lo, hi)`
will be
rearranged in such a way that the element at `mid` becomes the first
element of
the new range and the element at `mid - 1` becomes the last element,
all while
preserving the relative ordering of elements within the two
rotated subarrays.

All three versions below achieve the same result using in-place
algorithms.
Version 1 uses a straightforward swapping algorithm requiring
ForwardIterators.
Version 2 requires BidirectionalIterators, employing a well-known
trick with
three simple inversions. Version 3 requires RandomAccessIterators,
applying a
juggling algorithm which first divides the range into `gcd(hi - lo,
mid - lo)`
sets and then rotates the corresponding elements in each set.

Time Complexity:

- $O(n)$ per call to both functions, where n is the distance between lo and hi .

Space Complexity:

- $O(1)$ auxiliary.

*/

```
#include <algorithm>
```

```
template<class It>
```

```
void rotate1(It lo, It mid, It hi) {
```

```
    It next = mid;
```

```
    while (lo != next) {
```

```
        std::iter_swap(lo++, next++);
```

```
        if (next == hi) {
```

```
            next = mid;
```

```
        } else if (lo == mid) {
```

```
            mid = next;
```

```
        }
```

```
    }
```

```
}
```

```
template<class It>
```

```
void rotate2(It lo, It mid, It hi) {
```

```
    std::reverse(lo, mid);
```

```
    std::reverse(mid, hi);
```

```
    std::reverse(lo, hi);
```

```
}
```

```
int gcd(int a, int b) {
```

```
    return (b == 0) ? a : gcd(b, a % b);
```

```
}
```

```
template<class It>
```

```
void rotate3(It lo, It mid, It hi) {
```

```
    int n = hi - lo, jump = mid - lo;
```

```
    int g = gcd(jump, n), cycle = n / g;
```

```
    for (int i = 0; i < g; i++) {
```

```
        int curr = i, next;
```

```
        for (int j = 0; j < cycle - 1; j++) {
```

```
            next = curr + jump;
```

```
            if (next >= n) {
```

```
                next -= n;
```

```
            }
```

```
            std::iter_swap(lo + curr, lo + next);
```

```
            curr = next;
```

```
        }
```

```
    }
```

```
}
```

```
/** Example Usage and Output:
```

```

before sort:  2 4 2 0 5 10 7 3 7 1
after sort:   0 1 2 2 3 4 5 7 7 10
rotate left:  1 2 2 3 4 5 7 7 10 0
rotate right: 0 1 2 2 3 4 5 7 7 10

```

```

*** /

```

```

#include <algorithm>
#include <cassert>
#include <iostream>
#include <vector>
using namespace std;

```

```

int main() {
    vector<int> v0, v1, v2, v3;
    for (int i = 0; i < 10000; i++) {
        v0.push_back(i);
    }
    v1 = v2 = v3 = v0;
    int mid = 5678;
    std::rotate(v0.begin(), v0.begin() + mid, v0.end());
    rotatel(v1.begin(), v1.begin() + mid, v1.end());
    rotate2(v2.begin(), v2.begin() + mid, v2.end());
    rotate3(v3.begin(), v3.begin() + mid, v3.end());
    assert(v0 == v1 && v0 == v2 && v0 == v3);
}

```

```

// Example from:

```

<http://en.cppreference.com/w/cpp/algorithm/rotate>

```

int a[] = {2, 4, 2, 0, 5, 10, 7, 3, 7, 1};
vector<int> v(a, a + 10);
cout << "before sort: ";
for (int i = 0; i < (int)v.size(); i++) {
    cout << v[i] << " ";
}
cout << endl;

```

```

// Insertion sort.

```

```

for (vector<int>::iterator i = v.begin(); i != v.end(); ++i) {
    rotatel(std::upper_bound(v.begin(), i, *i), i, i + 1);
}
cout << "after sort: ";
for (int i = 0; i < (int)v.size(); i++) {
    cout << v[i] << " ";
}
cout << endl;

```

```

// Simple rotation to the left.

```

```

rotate2(v.begin(), v.begin() + 1, v.end());
cout << "rotate left: ";
for (int i = 0; i < (int)v.size(); i++) {

```

```

        cout << v[i] << " ";
    }
    cout << endl;

    // Simple rotation to the right.
    rotate3(v.rbegin(), v.rbegin() + 1, v.rend());
    cout << "rotate right: ";
    for (int i = 0; i < (int)v.size(); i++) {
        cout << v[i] << " ";
    }
    cout << endl;

    return 0;
}

```

Counting Inversions

```

/*
The number of inversions for an array a[] is defined as the number
of ordered
pairs (i, j) such that i < j and a[i] > a[j]. This is roughly how
"close" an
array is to being sorted, but is *not* the minimum number of swaps
required to
sort the array. If the array is sorted, then the inversion count
is 0. If the
array is sorted in decreasing order, then the inversion count is
maximal. The
following two functions are each techniques to efficiently count
inversions.
*/

```

```

#include <algorithm>
#include <iterator>
#include <vector>

```

```

/*
Version 1: Merge Sort

```

Returns the number of inversions given two RandomAccessIterators as a range [lo, hi). The range will become sorted after the function call. This requires operator < to be defined on the iterator's value type.

Time Complexity:

- $O(n \log n)$ per call to both functions, where n is distance between lo and hi in the first version and the number of array elements in the second version.

Space Complexity:

- $O(n)$ auxiliary heap space for both functions.

*/

```
template<class It>
long long inversions(It lo, It hi) {
    if (hi - lo < 2) {
        return 0;
    }
    It mid = lo + (hi - lo - 1)/2, a = lo, c = mid + 1;
    long long res = 0;
    res += inversions(lo, mid + 1);
    res += inversions(mid + 1, hi);
    typedef typename std::iterator_traits<It>::value_type T;
    std::vector<T> merged;
    while (a <= mid && c < hi) {
        if (*c < *a) {
            merged.push_back(*(c++));
            res += (mid - a) + 1;
        } else {
            merged.push_back(*(a++));
        }
    }
    if (a > mid) {
        for (It it = c; it != hi; ++it) {
            merged.push_back(*it);
        }
    } else {
        for (It it = a; it <= mid; ++it) {
            merged.push_back(*it);
        }
    }
    for (It it = lo; it != hi; ++it) {
        *it = merged[it - lo];
    }
    return res;
}
```

/*

Version 2: Power-of-Two Trick

Returns the number of inversions for an array `a[]` of `n` nonnegative integers.

After calling the function, every value of `a[]` will be set to 0.

Here, the time and space complexities depend on the magnitude of the maximum

value in `a[]`. Therefore for a running time of $O(n \log n)$, coordinate compression

may be applied to `a[]` so its maximum is strictly less than the length `n` itself.

Time Complexity: $O(m \log m)$, where m is maximum value in the array.

Space Complexity: $O(m)$ auxiliary.

*/

```
long long inversions(int n, int a[]) {
    int mx = 0;
    for (int i = 0; i < n; i++) {
        mx = std::max(mx, a[i]);
    }
    long long res = 0;
    std::vector<int> count(mx);
    while (mx > 0) {
        std::fill(count.begin(), count.end(), 0);
        for (int i = 0; i < n; i++) {
            if (a[i] % 2 == 0) {
                res += count[a[i] / 2];
            } else {
                count[a[i] / 2]++;
            }
        }
        mx = 0;
        for (int i = 0; i < n; i++) {
            mx = std::max(mx, a[i] /= 2);
        }
    }
    return res;
}
```

/** Example Usage **/

```
#include <cassert>
```

```
int main() {
    {
        int a[] = {6, 9, 1, 14, 8, 12, 3, 2};
        assert(inversions(a, a + 8) == 16);
    }
    {
        int a[] = {6, 9, 1, 14, 8, 12, 3, 2};
        assert(inversions(8, a) == 16);
    }
    return 0;
}
```

Coordinate Compression

/*

Given two ForwardIterators as a range [lo, hi) of n numerical elements, reassign each element in the range to an integer in $[0, k)$, where k is the

number of
distinct elements in the original range, while preserving the
initial relative
ordering of elements. That is, if `a[]` is an array of the original
values and `b[]`
is the compressed values, then every pair of indices `i, j` ($0 \leq i, j < n$) shall
satisfy `a[i] < a[j]` if and only if `b[i] < b[j]`.

Both implementations below require operator `<` to be defined on the
iterator's
value type. Version 1 performs the compression by sorting the
array, removing
duplicates, and binary searching for the position of each original
value.
Version 2 achieves the same result by inserting all values in a
balanced binary
search tree (`std::map`) which automatically removes duplicate
values and supports
efficient lookups of the compressed values.

Time Complexity:

- $O(n \log n)$ per call to either function, where `n` is the distance
between `lo` and
`hi`.

Space Complexity:

- $O(n)$ auxiliary heap space.

*/

```
#include <algorithm>
#include <iterator>
#include <map>
#include <vector>
```

```
template<class It> void compress1(It lo, It hi) {
    typedef typename std::iterator_traits<It>::value_type T;
    std::vector<T> v;
    for (It it = lo; it != hi; ++it) {
        v.push_back(*it);
    }
    std::sort(v.begin(), v.end());
    v.resize(std::unique(v.begin(), v.end()) - v.begin());
    for (It it = lo; it != hi; ++it) {
        *it = (int)(std::lower_bound(v.begin(), v.end(), *it) -
v.begin());
    }
}
```

```
template<class It> void compress2(It lo, It hi) {
    typedef typename std::iterator_traits<It>::value_type T;
```

```

std::map<T, int> m;
for (It it = lo; it != hi; ++it) {
    m[*it] = 0;
}
typename std::map<T, int>::iterator it = m.begin();
for (int id = 0; it != m.end(); it++) {
    it->second = id++;
}
for (It it = lo; it != hi; ++it) {
    *it = m[*it];
}
}

/** Example Usage and Output:

0 4 4 1 3 2 5 5
0 4 4 1 3 2 5 5
1 0 2 0 3 1

***/

#include <iostream>
using namespace std;

template<class It> void print_range(It lo, It hi) {
    while (lo != hi) {
        cout << *lo++ << " ";
    }
    cout << endl;
}

int main() {
    {
        int a[] = {1, 30, 30, 7, 9, 8, 99, 99};
        compress1(a, a + 8);
        print_range(a, a + 8);
    }
    {
        int a[] = {1, 30, 30, 7, 9, 8, 99, 99};
        compress2(a, a + 8);
        print_range(a, a + 8);
    }
    { // Non-integral types work too, as long as ints can be
      assigned to them.
        double a[] = {0.5, -1.0, 3, -1.0, 20, 0.5};
        compress1(a, a + 6);
        print_range(a, a + 6);
    }
    return 0;
}

```



```

Selection(Quickselect)
/*
nth_element2() is equivalent to std::nth_element(), taking
RandomAccessIterators
lo, nth, and hi as the range [lo, hi) to be partially sorted. The
values in
[lo, hi) are rearranged such that the value pointed to by nth is
the element
that would be there if the range were sorted. Furthermore, the
range is
partitioned such that no value in [lo, nth) compares greater than
the value
pointed to by nth and no value in (nth, hi) compares less. This
implementation
requires operator < to be defined on the iterator's value type.

```

Time Complexity:

```

- O(n) on average per call to nth_element2(), where n is the
distance between lo
and hi.

```

Space Complexity:

```

- O(1) auxiliary.

```

```

*/

```

```

#include <algorithm>

```

```

#include <cstdlib>

```

```

#include <iterator>

```

```

int rand32() {
    return (rand() & 0x7fff) | ((rand() & 0x7fff) << 15);
}

```

```

template<class It>

```

```

void nth_element2(It lo, It nth, It hi) {
    for (;;) {
        std::iter_swap(lo + rand32() % (hi - lo), hi - 1);
        typename std::iterator_traits<It>::value_type mid = *(hi - 1);
        It k = lo - 1;
        for (It it = lo; it != hi; ++it) {
            if (!(mid < *it)) {
                std::iter_swap(++k, it);
            }
        }
        if (nth < k) {
            hi = k;
        } else if (k < nth) {
            lo = k + 1;
        } else {
            return;
        }
    }
}

```

```

    }
}

/** Example Usage and Output:

2 3 3 4 5 6 6 7 9

***/

#include <cassert>
#include <iostream>
using namespace std;

template<class It>
void print_range(It lo, It hi) {
    while (lo != hi) {
        cout << *lo++ << " ";
    }
    cout << endl;
}

int main () {
    int n = 9;
    int a[] = {5, 6, 4, 3, 2, 6, 7, 9, 3};
    nth_element2(a, a + n/2, a + n);
    assert(a[n/2] == 5);
    print_range(a, a + n);
    return 0;
}

```

Longest Increasing Subsequence

/*
 Given two RandomAccessIterators lo and hi specifying a range [lo, hi), determine a longest subsequence of the range such that all of its elements are in strictly ascending order. This implementation requires operator < to be defined on the iterator's value type. The subsequence is not necessarily contiguous or unique, so only one such answer will be found. The answer is computed using binary search and dynamic programming.

Time Complexity:

- $O(n \log n)$ per call to `longest_increasing_subsequence()`, where n is the distance between `lo` and `hi`.

Space Complexity:

- $O(n)$ auxiliary heap space for `longest_increasing_subsequence()`.

```

*/

#include <iterator>
#include <vector>

template<class It>
std::vector<typename std::iterator_traits<It>::value_type>
longest_increasing_subsequence(It lo, It hi) {
    int len = 0, n = hi - lo;
    std::vector<int> prev(n), tail(n);
    for (int i = 0; i < n; i++) {
        int l = -1, h = len;
        while (h - l > 1) {
            int mid = (l + h)/2;
            if (*(lo + tail[mid]) < *(lo + i)) {
                l = mid;
            } else {
                h = mid;
            }
        }
        if (len < h + 1) {
            len = h + 1;
        }
        prev[i] = h > 0 ? tail[h - 1] : -1;
        tail[h] = i;
    }
    std::vector<typename std::iterator_traits<It>::value_type>
res(len);
    for (int i = tail[len - 1]; i != -1; i = prev[i]) {
        res[--len] = *(lo + i);
    }
    return res;
}

```

/** Example Usage and Output:

-5 1 9 10 11 13

*/

```

#include <iostream>
using namespace std;

template<class It> void print_range(It lo, It hi) {
    while (lo != hi) {
        cout << *lo++ << " ";
    }
    cout << endl;
}

```

```

int main () {

```

```

    int a[] = {-2, -5, 1, 9, 10, 8, 11, 10, 13, 11};
    vector<int> res = longest_increasing_subsequence(a, a + 10);
    print_range(res.begin(), res.end());
    return 0;
}

```

Maximal Subarray Sum (Kadane's)

```
/*
```

Given an array of numbers (at least one of which must be positive), determine the maximum possible sum of any contiguous subarray. Kadane's algorithm scans through the array, at each index computing the maximum positive sum subarray ending there. Either this subarray is empty (in which case its sum is zero) or it consists of one more element than the maximum sequence ending at the previous position. This can be adapted to compute the maximal submatrix sum as well.

```
*/
```

```

#include <algorithm>
#include <cstdlib>
#include <iterator>
#include <limits>
#include <vector>

```

```
/*
```

Returns the maximal subarray sum for the range [lo, hi), where lo and hi are RandomAccessIterators to numerical types. This implementation requires operators + and < to be defined on the iterator's value type. Optionally, two int pointers may be passed to store the inclusive boundary indices [res_lo, res_hi] of the resulting subarray. By convention, an input range consisting of only negative values will yield a size 1 subarray consisting of the maximum value.

Time Complexity:

- O(n) per call to max_subarray_sum(), where n is the distance between lo and hi.

Space Complexity:

- O(1) auxiliary.

*/

```
template<class It>
typename std::iterator_traits<It>::value_type
max_subarray_sum(It lo, It hi, int *res_lo = NULL, int *res_hi =
NULL) {
    typedef typename std::iterator_traits<It>::value_type T;
    int curr_begin = 0, begin = 0, end = -1;
    T sum = 0, max_sum = std::numeric_limits<T>::min();
    for (It it = lo; it != hi; ++it) {
        sum += *it;
        if (sum < 0) {
            sum = 0;
            curr_begin = (it - lo) + 1;
        } else if (max_sum < sum) {
            max_sum = sum;
            begin = curr_begin;
            end = it - lo;
        }
    }
    if (end == -1) {
        // All values negative. By convention, just return the maximum
        value.
        for (It it = lo; it != hi; ++it) {
            if (max_sum < *it) {
                max_sum = *it;
                begin = it - lo;
                end = begin;
            }
        }
    }
    if (res_lo != NULL && res_hi != NULL) {
        *res_lo = begin;
        *res_hi = end;
    }
    return max_sum;
}

/*
```

Returns the largest sum of any rectangular submatrix for a matrix of n rows by m columns. The matrix should be given as a 2-dimensional vector, where the outer vector must contain n vectors each of size m. This implementation requires operators + and < to be defined on the iterator's value type. Optionally, four int pointers may be passed to store the boundary indices of the

resulting subarray, with (r1, c1) specifying the top-left index and (r2, c2) specifying the bottom-right index. By convention, an input matrix consisting of only negative values will yield a size 1 submatrix consisting of the maximum value.

Time Complexity:

- $O(n*m^2)$ per call to `max_submatrix_sum()`, where n is the number of rows and m is the number of columns in the matrix.

Space Complexity:

- $O(n)$ auxiliary heap space for `max_submatrix_sum()`, where n is the number of rows in the matrix.

*/

```
template<class T>
T max_submatrix_sum(const std::vector<std::vector<T> > &matrix,
    int *r1 = NULL, int *c1 = NULL, int *r2 = NULL, int *c2 =
NULL) {
    int n = matrix.size(), m = matrix[0].size();
    std::vector<T> sums(n);
    T sum, max_sum = std::numeric_limits<T>::min();
    for (int clo = 0; clo < m; clo++) {
        std::fill(sums.begin(), sums.end(), 0);
        for (int chi = clo; chi < m; chi++) {
            for (int i = 0; i < n; i++) {
                sums[i] += matrix[i][chi];
            }
            int rlo, rhi;
            sum = max_subarray_sum(sums.begin(), sums.end(), &rlo,
&rhi);
            if (max_sum < sum) {
                max_sum = sum;
                if (r1 != NULL && c1 != NULL && r2 != NULL && c2 != NULL)
                {
                    *r1 = rlo;
                    *c1 = clo;
                    *r2 = rhi;
                    *c2 = chi;
                }
            }
        }
    }
    return max_sum;
}
```

/** Example Usage and Output:

Maximal sum subarray:

4 -1 2 1

Maximal sum submatrix:

9 2

-4 1

-1 8

*/

#include <cassert>

#include <iostream>

using namespace std;

int main() {

{

int a[] = {-2, -1, -3, 4, -1, 2, 1, -5, 4};

int lo = 0, hi = 0;

assert(max_subarray_sum(a, a + 3) == -1);

assert(max_subarray_sum(a, a + 9, &lo, &hi) == 6);

cout << "Maximal sum subarray:" << endl;

for (int i = lo; i <= hi; i++) {

cout << a[i] << " ";

}

cout << endl;

}

{

const int n = 4, m = 5;

int a[n][m] = {{0, -2, -7, 0, 5},

{9, 2, -6, 2, -4},

{-4, 1, -4, 1, 0},

{-1, 8, 0, -2, 3}};

vector<vector<int>> matrix(n);

for (int i = 0; i < n; i++) {

matrix[i] = vector<int>(a[i], a[i] + m);

}

int r1 = 0, c1 = 0, r2 = 0, c2 = 0;

assert(max_submatrix_sum(matrix, &r1, &c1, &r2, &c2) == 15);

cout << "\nMaximal sum submatrix:" << endl;

for (int i = r1; i <= r2; i++) {

for (int j = c1; j <= c2; j++) {

cout << matrix[i][j] << " ";

}

cout << endl;

}

}

return 0;

}

Majority Element (Boyer-Moore)

/*

Given two ForwardIterators lo and hi specifying a range [lo, hi) of n elements, return an iterator to the first occurrence of the majority element, or the iterator hi if there is no majority element. The majority element is defined as an element which occurs strictly more than $\text{floor}(n/2)$ times in the range. This implementation requires operator == to be defined on the iterator's value type.

Time Complexity:

- $O(n)$ per call to majority(), where n is the size of the array.

Space Complexity:

- $O(1)$ auxiliary.

*/

```
template<class It>
It majority(It lo, It hi) {
    int count = 0;
    It candidate = lo;
    for (It it = lo; it != hi; ++it) {
        if (count == 0) {
            candidate = it;
            count = 1;
        } else if (*it == *candidate) {
            count++;
        } else {
            count--;
        }
    }
    count = 0;
    for (It it = lo; it != hi; ++it) {
        if (*it == *candidate) {
            count++;
        }
    }
    if (count <= (hi - lo)/2) {
        return hi;
    }
    return candidate;
}
```

/** Example Usage **/

#include <cassert>


```

int main() {
    int a[] = {3, 2, 3, 1, 3};
    assert(*majority(a, a + 5) == 3);
    int b[] = {2, 3, 3, 3, 2, 1};
    assert(majority(b, b + 6) == b + 6);
    return 0;
}

```

Subset Sum (Meet-in-the-Middle)
/*

Given RandomAccessIterators lo and hi specifying a range [lo, hi) of integers,
return the minimum sum of any subset of the range that is greater than or equal
to a given integer v. This is a generalization of the NP-complete subset sum
problem, which asks whether a subset summing to 0 exists (equivalent in this
case to checking if $v = 0$ yields an answer of 0). This implementation uses a
meet-in-the-middle algorithm to precompute and search for a lower bound. Note
that 64-bit integers are used in intermediate calculations to avoid overflow.

Time Complexity:

- $O(n \cdot 2^{(n/2)})$ per call to `sum_lower_bound()`, where n is the distance between `lo` and `hi`.

Space Complexity:

- $O(n)$ auxiliary heap space, where n is the number of array elements.

*/

```

#include <algorithm>
#include <limits>
#include <vector>

```

```

template<class It>
long long sum_lower_bound(It lo, It hi, long long v) {
    int n = hi - lo, llen = 1 << (n/2), hlen = 1 << (n - n/2);
    std::vector<long long> lsum(llen), hsum(hlen);
    for (int mask = 0; mask < llen; mask++) {
        for (int i = 0; i < n/2; i++) {
            if ((mask >> i) & 1) {
                lsum[mask] += *(lo + i);
            }
        }
    }
}

```

```

    }
}
for (int mask = 0; mask < hlen; mask++) {
    for (int i = 0; i < (n - n/2); i++) {
        if ((mask >> i) & 1) {
            hsum[mask] += *(lo + i + n/2);
        }
    }
}
std::sort(lsum.begin(), lsum.end());
std::sort(hsum.begin(), hsum.end());
int l = 0, h = hlen - 1;
long long curr = std::numeric_limits<long long>::min();
while (l < llen && h >= 0) {
    if (lsum[l] + hsum[h] <= v) {
        curr = std::max(curr, lsum[l] + hsum[h]);
        l++;
    } else {
        h--;
    }
}
return curr;
}

```

/** Example Usage */

```
#include <cassert>
```

```

int main() {
    int a[] = {9, 1, 5, 0, 1, 11, 5};
    assert(sum_lower_bound(a, a + 7, 8) == 7);
    int b[] = {-7, -3, -2, 5, 8};
    assert(sum_lower_bound(b, b + 5, 0) == 0);
    return 0;
}

```

Maximal Zero Submatrix

/*

Given a rectangular matrix with n rows and m columns consisting of only 0's and 1's as a two-dimensional vector of bool, return the area of the largest rectangular submatrix consisting of only 0's. This solution uses a reduction to the problem of finding the maximum rectangular area under a histogram, which is efficiently solved using a stack algorithm.

Time Complexity:

- $O(n*m)$ per call to `max_zero_submatrix()`, where n is the number of rows and m is the number of columns in the matrix.

Space Complexity:

- $O(m)$ auxiliary heap space, where m is the number of columns in the matrix.

*/

```
#include <algorithm>
#include <stack>
#include <vector>
```

```
int max_zero_submatrix(const std::vector<std::vector<bool> >
&matrix) {
    int n = matrix.size(), m = matrix[0].size(), res = 0;
    std::vector<int> d(m, -1), d1(m), d2(m);
    for (int r = 0; r < n; r++) {
        for (int c = 0; c < m; c++) {
            if (matrix[r][c]) {
                d[c] = r;
            }
        }
        std::stack<int> s;
        for (int c = 0; c < m; c++) {
            while (!s.empty() && d[s.top()] <= d[c]) {
                s.pop();
            }
            d1[c] = s.empty() ? -1 : s.top();
            s.push(c);
        }
        while (!s.empty()) {
            s.pop();
        }
        for (int c = m - 1; c >= 0; c--) {
            while (!s.empty() && d[s.top()] <= d[c]) {
                s.pop();
            }
            d2[c] = s.empty() ? m : s.top();
            s.push(c);
        }
        for (int j = 0; j < m; j++) {
            res = std::max(res, (r - d[j])*(d2[j] - d1[j] - 1));
        }
    }
    return res;
}
```

/** Example Usage **/

```

#include <cassert>
using namespace std;

int main() {
    const int n = 5, m = 6;
    bool a[n][m] = {{1, 0, 1, 1, 0, 0},
                    {1, 0, 0, 1, 0, 0},
                    {0, 0, 0, 0, 0, 1},
                    {1, 0, 0, 1, 0, 0},
                    {1, 0, 1, 0, 0, 1}};
    vector<vector<bool> > matrix(n);
    for (int i = 0; i < n; i++) {
        matrix[i] = vector<bool>(a[i], a[i] + m);
    }
    assert(max_zero_submatrix(matrix) == 6);
    return 0;
}

```

Binary Search

/*

Binary search can be generally used to find the input value corresponding to any output value of a monotonic (strictly non-increasing or strictly non-decreasing) function in $O(\log n)$ time with respect to the domain size. This is a special case of finding the exact point at which any given monotonic Boolean function changes from true to false or vice versa. Unlike searching through an array, discrete binary search is not restricted by available memory, making it useful for handling infinitely large search spaces such as real number intervals.

`binary_search_first_true()` takes two integers `lo` and `hi` as boundaries for the search space `[lo, hi)` (i.e. including `lo`, but excluding `hi`) and returns the smallest integer `k` in `[lo, hi)` for which the predicate `pred(k)` tests true. If `pred(k)` tests false for every `k` in `[lo, hi)`, then `hi` is returned. This function must be used on a range in which there exists a constant `k` such that `pred(x)` tests false for every `x` in `[lo, k)` and true for every `x` in `[k, hi)`.

`binary_search_last_true()` takes two integers `lo` and `hi` as

boundaries for the search space $[lo, hi)$ (i.e. including lo , but excluding hi) and returns the largest integer k in $[lo, hi)$ for which the predicate $pred(k)$ tests true. If $pred(k)$ tests false for every k in $[lo, hi)$, then hi is returned. This function must be used on a range in which there exists a constant k such that $pred(x)$ tests true for every x in $[lo, k]$ and false for every x in (k, hi) .

Time Complexity:

- $O(\log n)$ calls will be made to $pred()$ in either function, where n is the distance between lo and hi .

Space Complexity:

- $O(1)$ auxiliary.

*/

```
template<class Int, class IntPredicate>
Int binary_search_first_true(Int lo, Int hi, IntPredicate pred)
{ // 000[1]11
    Int mid, _hi = hi;
    while (lo < hi) {
        mid = lo + (hi - lo)/2;
        if (pred(mid)) {
            hi = mid;
        } else {
            lo = mid + 1;
        }
    }
    if (!pred(lo)) {
        return _hi; // All false.
    }
    return lo;
}
```

```
template<class Int, class IntPredicate>
Int binary_search_last_true(Int lo, Int hi, IntPredicate pred)
{ // 11[1]000
    Int mid, _hi = hi--;
    while (lo < hi) {
        mid = lo + (hi - lo + 1)/2;
        if (pred(mid)) {
            lo = mid;
        } else {
            hi = mid - 1;
        }
    }
}
```

```

    }
    if (!pred(lo)) {
        return _hi; // All false.
    }
    return lo;
}

```

/*

fbinary_search() is the equivalent of binary_search_first_true() on floating point predicates. Since any interval of real numbers is dense, the exact target cannot be found due to floating point error. Instead, the function returns a value that is very close to the border between false and true. The precision of the answer depends on the number of repetitions the function performs. Since each repetition bisects the search space, the absolute error of the answer is $1/(2^r)$ times the distance between lo and hi after r repetitions. Although it is possible to control the error by looping while hi - lo is greater than an arbitrary epsilon, it is simpler to let the loop run for a desired number of iterations until floating point arithmetic break down. 100 iterations is usually sufficient, since the search space will be reduced to 2^{-100} (roughly 10^{-30}) times its original size.

This implementation can be modified to find the "last true" point in the range by simply interchanging the assignments of lo and hi in the if-else statements.

Time Complexity:

- $O(\log n)$ calls will be made to pred(), where n is the distance between lo and hi divided by the desired absolute error (based on the number of iterations).

Space Complexity:

- $O(1)$ auxiliary.

*/

```

template<class DoublePredicate>
double fbinary_search(double lo, double hi, DoublePredicate pred)

```

```

{ // 000[1]11
  double mid;
  for (int i = 0; i < 100; i++) {
    mid = (lo + hi)/2.0;
    if (pred(mid)) {
      hi = mid;
    } else {
      lo = mid;
    }
  }
  return lo;
}

/** Example Usage */

#include <cassert>
#include <cmath>

// Simple predicate examples.
bool pred1(int x) { return x >= 3; }
bool pred2(int x) { return false; }
bool pred3(int x) { return x <= 5; }
bool pred4(int x) { return true; }
bool pred5(double x) { return x >= 1.2345; }

int main() {
  assert(binary_search_first_true(0, 7, pred1) == 3);
  assert(binary_search_first_true(0, 7, pred2) == 7);
  assert(binary_search_last_true(0, 7, pred3) == 5);
  assert(binary_search_last_true(0, 7, pred4) == 6);
  assert(fabs(fbinary_search(-10.0, 10.0, pred5) - 1.2345) < 1e-
15);
  return 0;
}

```

Ternary Search

/*

Given a unimodal function $f(x)$ taking a single double argument,
find its global
maximum or minimum point to a specified absolute error.

`ternary_search_min()` takes the domain $[lo, hi]$ of a continuous
function $f(x)$ and
returns a number x such that f is strictly decreasing on the
interval $[lo, x]$
and strictly increasing on the interval $[x, hi]$. For the function
to be correct
and deterministic, such an x must exist and be unique.

ternary_search_max() takes the domain [lo, hi] of a continuous function $f(x)$ and returns a number x such that f is strictly increasing on the interval [lo, x] and strictly decreasing on the interval [x , hi]. For the function to be correct and deterministic, such an x must exist and be unique.

Time Complexity:

- $O(\log n)$ calls will be made to $f()$, where n is the distance between lo and hi divided by the specified absolute error (epsilon).

Space Complexity:

- $O(1)$ auxiliary.

*/

```
template<class UnimodalFunction>
double ternary_search_min(double lo, double hi, UnimodalFunction
f) {
    static const double EPS = 1e-12;
    double lthird, hthird;
    while (hi - lo > EPS) {
        lthird = lo + (hi - lo)/3;
        hthird = hi - (hi - lo)/3;
        if (f(lthird) < f(hthird)) {
            hi = hthird;
        } else {
            lo = lthird;
        }
    }
    return lo;
}
```

```
template<class UnimodalFunction>
double ternary_search_max(double lo, double hi, UnimodalFunction
f) {
    static const double EPS = 1e-12;
    double lthird, hthird;
    while (hi - lo > EPS) {
        lthird = lo + (hi - lo)/3;
        hthird = hi - (hi - lo)/3;
        if (f(lthird) < f(hthird)) {
            lo = lthird;
        } else {
            hi = hthird;
        }
    }
    return hi;
}
```



```

/** Example Usage */

#include <cassert>
#include <cmath>

bool equal(double a, double b) {
    return fabs(a - b) < 1e-7;
}

// Parabola opening up with vertex at (-2, -24).
double f1(double x) {
    return 3*x*x + 12*x - 12;
}

// Parabola opening down with vertex at (2/19, 8366/95) ~ (0.105,
88.063).
double f2(double x) {
    return -5.7*x*x + 1.2*x + 88;
}

// Absolute value function shifted to the right by 30 units.
double f3(double x) {
    return fabs(x - 30);
}

int main() {
    assert(equal(ternary_search_min(-1000, 1000, f1), -2));
    assert(equal(ternary_search_max(-1000, 1000, f2), 2.0/19));
    assert(equal(ternary_search_min(-1000, 1000, f3), 30));
    return 0;
}

```

Hill Climbing

/*

Given a continuous function $f(x, y)$ to double and a (possibly arbitrary) initial guess (x_0, y_0) , return a potential global minimum found through hill-climbing. Optionally, two double pointers `critical_x` and `critical_y` may be passed to store the input points to f at which the returned minimum value is attained.

Hill-climbing is a heuristic which starts at the guess, then considers taking a single step in each of a fixed number of directions. The direction with the best (in this case, minimum) value is chosen, and further steps

are taken in it
until the answer no longer improves. When this happens, the step
size is reduced
and the same process repeats until a desired absolute error is
reached. The
technique's success heavily depends on the behavior of f and the
initial guess.
Therefore, the result is not guaranteed to be the global minimum.

Time Complexity:

- $O(d \log n)$ call will be made to f , where d is the number of directions considered at each position and n is the search space that is approximately proportional to the maximum possible step size divided by the minimum possible step size.

Space Complexity:

- $O(1)$ auxiliary.

*/

```
#include <cstdlib>
```

```
#include <cmath>
```

```
template<class ContinuousFunction>
```

```
double find_min(ContinuousFunction f, double x0, double y0,  
                double *critical_x = NULL, double *critical_y =  
NULL) {
```

```
    static const double PI = acos(-1.0);  
    static const double STEP_MIN = 1e-9, STEP_MAX = 1e6;  
    static const int NUM_DIRECTIONS = 6;  
    double x = x0, y = y0, res = f(x0, y0);  
    for (double step = STEP_MAX; step > STEP_MIN; ) {  
        double best = res, best_x = x, best_y = y;  
        bool found = false;  
        for (int i = 0; i < NUM_DIRECTIONS; i++) {  
            double a = 2.0*PI*i / NUM_DIRECTIONS;  
            double x2 = x + step*cos(a), y2 = y + step*sin(a);  
            double value = f(x2, y2);  
            if (best > value) {  
                best_x = x2;  
                best_y = y2;  
                best = value;  
                found = true;  
            }  
        }  
        if (!found) {  
            step /= 2.0;  
        } else {
```

```

        x = best_x;
        y = best_y;
        res = best;
    }
}
if (critical_x != NULL && critical_y != NULL) {
    *critical_x = x;
    *critical_y = y;
}
return res;
}

```

/** Example Usage **/

```

#include <cassert>
#include <cmath>

```

```

bool eq(double a, double b) {
    return fabs(a - b) < 1e-8;
}

```

```

// Paraboloid with global minimum at f(2, 3) = 0.
double f(double x, double y) {
    return (x - 2)*(x - 2) + (y - 3)*(y - 3);
}

```

```

int main() {
    double x, y;
    assert(eq(find_min(f, 0, 0, &x, &y), 0));
    assert(eq(x, 2) && eq(y, 3));
    return 0;
}

```

Convex Hull Trick (Semi-Dynamic)

/*

Given a set of pairs (m, b) specifying lines of the form $y = mx + b$, process a set of x-coordinate queries each asking to find the minimum y-value when any of the given lines are evaluated at the specified x. For each `add_line(m, b)` call, m must be the minimum m of all lines added so far. For each `query(x)` call, x must be the maximum x of all queries made so far.

The following implementation is a concise, semi-dynamic version of the convex hull optimization technique. It supports an an interlaced sequence of `add_line()`

and query() calls, as long as the preconditions of descending m and ascending x are satisfied. As a result, it may be necessary to sort the lines and queries before calling the functions. In that case, the overall time complexity will be dominated by the sorting step.

Time Complexity:

- O(n) for any interlaced sequence of add_line() and query() calls, where n is the number of lines added. This is because the overall number of steps taken by add_line() and query() are respectively bounded by the number of lines.
- Thus a single call to either add_line() or query() will have an amortized O(1) running time.

Space Complexity:

- O(n) for storage of the lines.
- O(1) auxiliary for add_line() and query().

*/

```
#include <vector>
```

```
std::vector<long long> M, B;
int ptr = 0;
```

```
void add_line(long long m, long long b) {
    int len = M.size();
    while (len > 1 && (B[len - 2] - B[len - 1])*(m - M[len - 1]) >=
                (B[len - 1] - b)*(M[len - 1] - M[len - 2])) {
        len--;
    }
    M.resize(len);
    B.resize(len);
    M.push_back(m);
    B.push_back(b);
}
```

```
long long query(long long x) {
    if (ptr >= (int)M.size()) {
        ptr = (int)M.size() - 1;
    }
    while (ptr + 1 < (int)M.size() &&
            M[ptr + 1]*x + B[ptr + 1] <= M[ptr]*x + B[ptr]) {
        ptr++;
    }
    return M[ptr]*x + B[ptr];
}
```

```

}

/** Example Usage */

#include <cassert>

int main() {
    add_line(3, 0);
    add_line(2, 1);
    add_line(1, 2);
    add_line(0, 6);
    assert(query(0) == 0);
    assert(query(1) == 3);
    assert(query(2) == 4);
    assert(query(3) == 5);
    return 0;
}

```

Convex Hull Trick (Fully Dynamic)

/*

Given a set of pairs (m, b) specifying lines of the form $y = mx + b$, process a set of x-coordinate queries each asking to find the minimum y-value when any of the given lines are evaluated at the specified x. To instead have the queries optimize for maximum y-value, call the constructor with `query_max=true`.

The following implementation is a fully dynamic variant of the convex hull optimization technique, using a self-balancing binary search tree (`std::set`) to support the ability to call `add_line()` and `query()` in any desired order.

Time Complexity:

- $O(n)$ for any interlaced sequence of `add_line()` and `query()` calls, where n is the number of lines added. This is because the overall number of steps taken by `add_line()` and `query()` are respectively bounded by the number of lines. Thus a single call to either `add_line()` or `query()` will have an $O(1)$ amortized running time.

Space Complexity:

- $O(n)$ for storage of the lines.

```

- O(1) auxiliary for add_line() and query().

*/

#include <limits>
#include <set>

class hull_optimizer {
    struct line {
        long long m, b, value;
        double xlo;
        bool is_query, query_max;

        line(long long m, long long b, long long v, bool is_query,
            bool query_max)
            : m(m), b(b), value(v), xlo(-
std::numeric_limits<double>::max()),
            is_query(is_query), query_max(query_max) {}

        double intersect(const line &l) const {
            if (m == l.m) {
                return std::numeric_limits<double>::max();
            }
            return (double)(l.b - b)/(m - l.m);
        }

        bool operator<(const line &l) const {
            if (l.is_query) {
                return query_max ? (xlo < l.value) : (l.value < xlo);
            }
            return m < l.m;
        }
    };

    std::set<line> hull;
    bool query_max;

    typedef std::set<line>::iterator hulliter;

    bool has_prev(hulliter it) const {
        return it != hull.begin();
    }

    bool has_next(hulliter it) const {
        return (it != hull.end()) && (++it != hull.end());
    }

    bool irrelevant(hulliter it) const {
        if (!has_prev(it) || !has_next(it)) {
            return false;
        }
    }
};

```

```

    hulliter prev = it, next = it;
    --prev;
    ++next;
    return query_max ? (prev->intersect(*next) <= prev-
>intersect(*it))
                        : (next->intersect(*prev) <= next-
>intersect(*it));
}

hulliter update_left_border(hulliter it) {
    if ((query_max && !has_prev(it)) || (!query_max
&& !has_next(it))) {
        return it;
    }
    hulliter it2 = it;
    double value = it->intersect(query_max ? *--it2 : *++it2);
    line l(*it);
    l.xlo = value;
    hull.erase(it++);
    return hull.insert(it, l);
}

public:
    hull_optimizer(bool query_max = false) : query_max(query_max) {}

    void add_line(long long m, long long b) {
        line l(m, b, 0, false, query_max);
        hulliter it = hull.lower_bound(l);
        if (it != hull.end() && it->m == l.m) {
            if ((query_max && it->b < b) || (!query_max && b < it->b)) {
                hull.erase(it++);
            } else {
                return;
            }
        }
        it = hull.insert(it, l);
        if (irrelevant(it)) {
            hull.erase(it);
            return;
        }
        while (has_prev(it) && irrelevant(--it)) {
            hull.erase(it++);
        }
        while (has_next(it) && irrelevant(++it)) {
            hull.erase(it--);
        }
        it = update_left_border(it);
        if (has_prev(it)) {
            update_left_border(--it);
        }
        if (has_next(++it)) {

```

```

        update_left_border(++it);
    }
}

long long query(long long x) const {
    line q(0, 0, x, true, query_max);
    hulliter it = hull.lower_bound(q);
    if (query_max) {
        --it;
    }
    return it->m*x + it->b;
}
};

```

/** Example Usage **/

```
#include <cassert>
```

```

int main() {
    hull_optimizer h;
    h.add_line(3, 0);
    h.add_line(0, 6);
    h.add_line(1, 2);
    h.add_line(2, 1);
    assert(h.query(0) == 0);
    assert(h.query(2) == 4);
    assert(h.query(1) == 3);
    assert(h.query(3) == 5);
    return 0;
}

```

Cycle Detection (Floyd's)

/*

Given a function f mapping a set of integers to itself and an x -coordinate in the set, return a pair containing the (position, length) of a cycle in the sequence of numbers obtained from repeatedly composing f with itself starting with the initial x . Formally, since f maps a finite set S to itself, some value is guaranteed to eventually repeat in the sequence:

$$x[0], x[1]=f(x[0]), x[2]=f(x[1]), \dots, x[n]=f(x[n-1]), \dots$$

There must exist a pair of indices i and j ($i < j$) such that $x[i] = x[j]$. When this happens, the rest of the sequence will consist of the subsequence from $x[i]$ to $x[j-1]$ repeating indefinitely. The cycle detection problem

asks to find
such an i , along with the length of the repeating subsequence. A
well-known
special case is the problem of cycle-detection in a degenerate
linked list.

Floyd's cycle-finding algorithm, a.k.a. the "tortoise and the hare
algorithm",
is a space-efficient algorithm that moves two pointers through the
sequence at
different speeds. Each step in the algorithm moves the "tortoise"
one step
forward and the "hare" two steps forward in the sequence,
comparing the sequence
values at each step. The first value which is simultaneously
pointed to by both
pointers is the start of the sequence.

Time Complexity:

- $O(m + n)$ per call to `find_cycle_floyd()`, where m is the smallest
index of the
sequence which is the beginning of a cycle, and n is the cycle's
length.

Space Complexity:

- $O(1)$ auxiliary.

*/

```
#include <utility>
```

```
template<class IntFunction>
std::pair<int, int> find_cycle_floyd(IntFunction f, int x0) {
    int tortoise = f(x0), hare = f(f(x0));
    while (tortoise != hare) {
        tortoise = f(tortoise);
        hare = f(f(hare));
    }
    int start = 0;
    tortoise = x0;
    while (tortoise != hare) {
        tortoise = f(tortoise);
        hare = f(hare);
        start++;
    }
    int length = 1;
    hare = f(tortoise);
    while (tortoise != hare) {
        hare = f(hare);
        length++;
    }
}
```

```

    return std::make_pair(start, length);
}

/** Example Usage */

#include <cassert>
#include <set>
using namespace std;

int f(int x) {
    return (123*x*x + 4567890) % 1337;
}

void verify(int x0, int start, int length) {
    set<int> s;
    int x = x0;
    for (int i = 0; i < start; i++) {
        assert(!s.count(x));
        s.insert(x);
        x = f(x);
    }
    int startx = x;
    s.clear();
    for (int i = 0; i < length; i++) {
        assert(!s.count(x));
        s.insert(x);
        x = f(x);
    }
    assert(startx == x);
}

int main () {
    int x0 = 0;
    pair<int, int> res = find_cycle_floyd(f, x0);
    assert(res == make_pair(4, 2));
    verify(x0, res.first, res.second);
    return 0;
}

```

Cycle Detection (Brent's)

/*

Given a function f mapping a set of integers to itself and an x -coordinate in the set, return a pair containing the (position, length) of a cycle in the sequence of numbers obtained from repeatedly composing f with itself starting with the initial x . Formally, since f maps a finite set S to itself, some value is guaranteed to eventually repeat in the sequence:

$x[0], x[1]=f(x[0]), x[2]=f(x[1]), \dots, x[n]=f(x[n-1]), \dots$

There must exist a pair of indices i and j ($i < j$) such that $x[i] = x[j]$. When this happens, the rest of the sequence will consist of the subsequence from $x[i]$ to $x[j-1]$ repeating indefinitely. The cycle detection problem asks to find such an i , along with the length of the repeating subsequence. A well-known special case is the problem of cycle-detection in a degenerate linked list.

While Floyd's cycle-finding algorithm finds cycles by simultaneously moving two pointers at different speeds, Brent's algorithm keeps the tortoise pointer stationary in every iteration, only teleporting it to the hare pointer at every power of two. The smallest power of two at which they meet is the start of the first cycle. This improves upon the constant factor of Floyd's algorithm by reducing the number of calls made to f .

Time Complexity:

- $O(m+n)$ per call to `find_cycle_brent()`, where m is the smallest index of the sequence which is the beginning of a cycle, and n is the cycle's length.

Space Complexity:

- $O(1)$ auxiliary.

*/

#include <utility>

```
template<class IntFunction>
std::pair<int,int> find_cycle_brent(IntFunction f,int x0){
    int power=1,length=1,tortoise=x0,hare=f(x0);
    while (tortoise!=hare){
        if (power==length){
            tortoise=hare;
            power*=2;
            length=0;
        }
        hare=f(hare);
        length++;
    }
    hare=x0;
    for (int i=0;i<length;i++){
        hare=f(hare);
```

```

    }
    int start = 0;
    tortoise = x0;
    while (tortoise != hare) {
        tortoise = f(tortoise);
        hare = f(hare);
        start++;
    }
    return std::make_pair(start, length);
}

/** Example Usage */

#include <cassert>
#include <set>
using namespace std;

int f(int x) {
    return (123 * x * x + 4567890) % 1337;
}

void verify(int x0, int start, int length) {
    set<int> s;
    int x = x0;
    for (int i = 0; i < start; i++) {
        assert(!s.count(x));
        s.insert(x);
        x = f(x);
    }
    int startx = x;
    s.clear();
    for (int i = 0; i < length; i++) {
        assert(!s.count(x));
        s.insert(x);
        x = f(x);
    }
    assert(startx == x);
}

int main() {
    int x0 = 0;
    pair<int, int> res = find_cycle_brent(f, x0);
    assert(res == make_pair(4, 2));
    verify(x0, res.first, res.second);
    return 0;
}

```

Binary Exponentiation

/*

Given three unsigned 64-bit integers x , n , and m , `powmod()` returns x raised to the power of n (mod m). `mulmod()` returns x multiplied by n (mod m). Despite the fact that both functions use unsigned 64-bit integers for arguments and intermediate calculations, arguments x and n must not exceed $2^{63} - 1$ (the maximum value of a signed 64-bit integer) for the result to be correctly computed without overflow.

Binary exponentiation, also known as exponentiation by squaring, decomposes the exponentiation into a logarithmic number of multiplications while avoiding overflow. To further prevent overflow in the intermediate squaring computations, multiplication is performed using a similar principle of repeated addition.

Time Complexity:

- $O(\log n)$ per call to `mulmod()` and `powmod()`, where n is the second argument.

Space Complexity:

- $O(1)$ auxiliary.

*/

```
typedef unsigned long long uint64;
```

```
uint64 mulmod(uint64 x, uint64 n, uint64 m) {
    uint64 a = 0, b = x % m;
    for (; n > 0; n >>= 1) {
        if (n & 1) {
            a = (a + b) % m;
        }
        b = (b << 1) % m;
    }
    return a % m;
}
```

```
uint64 powmod(uint64 x, uint64 n, uint64 m) {
    uint64 a = 1, b = x;
    for (; n > 0; n >>= 1) {
        if (n & 1) {
            a = mulmod(a, b, m);
        }
    }
}
```

```
    b = mulmod(b,b,m);
}
return a % m;
}

/** Example Usage */

#include <cassert>

int main(){
    assert(powmod(2,10,1000000007) == 1024);
    assert(powmod(2,62,1000000) == 387904);
    assert(powmod(10001,10001,100000) == 10001);
    return 0;
}
```